

# Answer correctness vs calibration success — Gemma 3 4B IT

Three deterministically-graded tools. Top-right cell = correct answer despite calibration failure; bottom-left = calibration ✓ but wrong answer (bad tool args, or tool call without prose answer).



